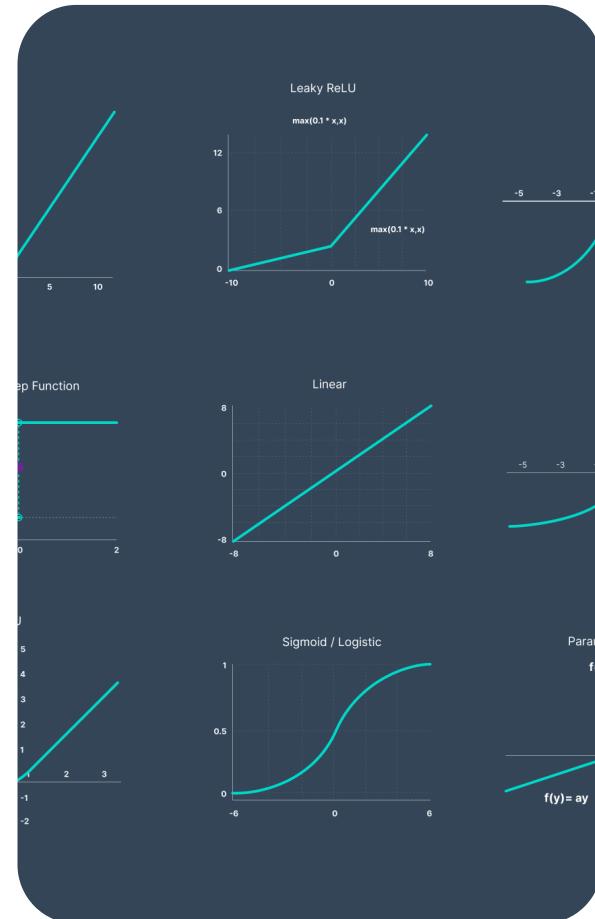


BLOG > DEEP LEARNING

Activation Functions in Neural Networks [12 Types & Use Cases]

What is a neural network activation function and how does it work? Explore twelve different types of activation functions and learn how to pick the right one.



🕒 14 min read · May 27, 2021

V7 Platform

Industries

Company

Jobs
20

Resources

Pricing

Log in

Request
a demo

What is
a
Neural

"*The world is one big data problem.*"
As it turns out—

A COMPREHENSIVE
GUIDE

🍪 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



Accept All

Manage cookies

Need
an
Activation
Function?

3 Types
of
Neural
Networks
Activation
Functions

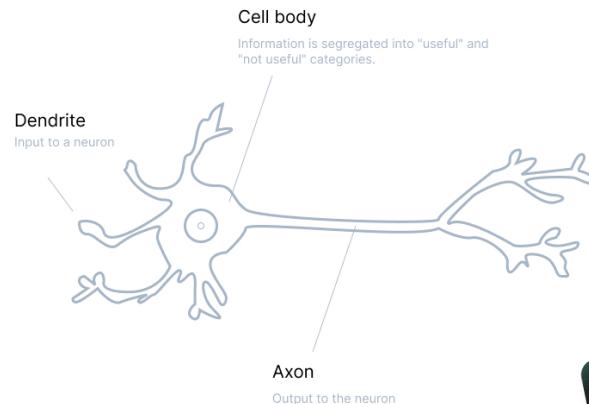
Why
are
deep
neural
networks
hard to
train?

How to
choose
the
right
Activation
Function?

Neural
Networks
Activation
Functions
in a
Nutshell

Text
Link

Every single moment our brain is trying to segregate the incoming information into the “useful” and “not-so-useful” categories.



We tackle considerations for building or buying an ML Ops platform, from data security, to costs and cutting-edge features.

[Download](#)

By submitting you are agreeing to V7's [privacy policy](#) and to receive other content from V7.



A similar process occurs in [artificial neural network architectures](#) in deep learning.

The segregation plays a key role in helping a neural network properly function, ensuring that it learns from the useful information rather than get stuck analyzing the not-useful part.

And this is also where activation functions come into the picture.

 **Activation Function** helps the neural network to use important information while suppressing irrelevant data points.

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.

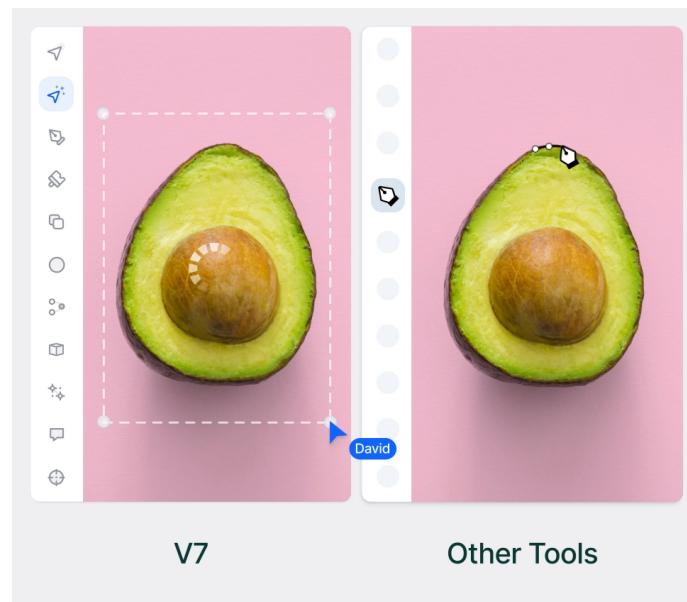
X

Here's what we'll cover:

1. [What is a Neural Network Activation Function?](#)
2. [Why do Neural Networks Need an Activation Function?](#)
3. [3 Types of Neural Networks Activation Functions](#)
4. [Why are deep neural networks hard to train?](#)
5. [How to choose the right Activation Function?](#)
6. [Neural Networks Activation Functions in a Nutshell](#)

Ready? Let's get started :)

Solve any video or image labeling task 10x faster and with 10x less manual work.



[Try V7 Now](#)

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



What is a Neural Network Activation Function?

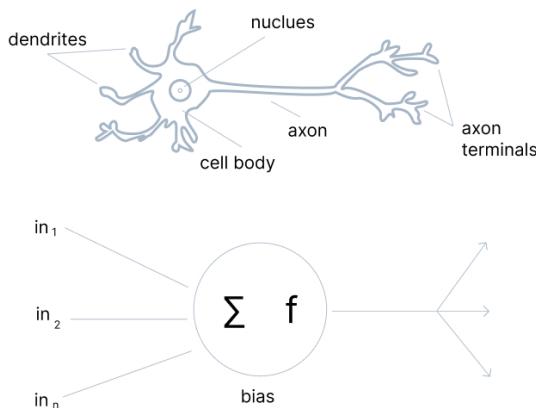
An **Activation Function** decides whether a neuron should be activated or not. This means that it will decide whether the neuron's input to the network is important or not in the process of prediction using simpler mathematical operations.

The role of the Activation Function is to derive output from a set of input values fed to a node (or a layer).

But—

Let's take a step back and clarify: What exactly is a *node*?

Well, if we compare the neural network to our brain, a node is a replica of a neuron that receives a set of input signals—external stimuli.



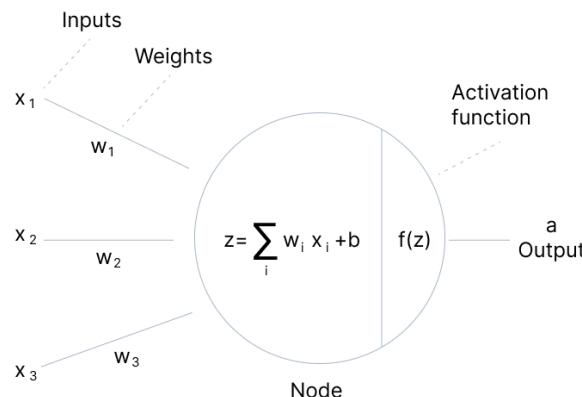
We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



Depending on the nature and intensity of these input signals, the brain processes them and decides whether the neuron should be activated (“fired”) or not.

In **deep learning**, this is also the role of the Activation Function—that’s why it’s often referred to as a *Transfer Function* in Artificial Neural Network.

The primary role of the Activation Function is to transform the summed weighted input from the node into an output value to be fed to the next hidden layer or as output.



V7 Labs

Now, let's have a look at the Neural Networks Architecture.

Elements of a Neural Networks Architecture

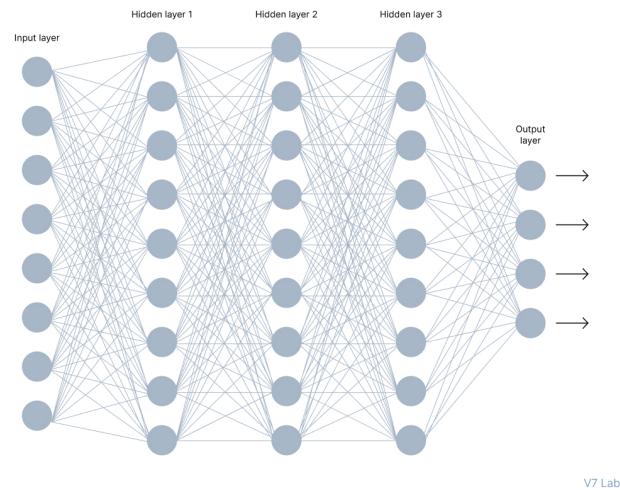
Here's the thing—

If you don't understand the concept of neural

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



structure of the Neural Networks Architecture and its components. Here it is.



V7 Labs

In the image above, you can see a neural network made of interconnected neurons. Each of them is characterized by its **weight**, **bias**, and **activation function**.

Here are other elements of this network.

Input Layer

The input layer takes raw input from the domain. No computation is performed at this layer. Nodes here just pass on the information (features) to the hidden layer.

Hidden Layer

As the name suggests, the nodes of this layer are not exposed. They provide an abstraction to the neural network.

The hidden layer performs all kinds of computation on the features entered through

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



It's the final layer of the network that brings the information learned through the hidden layer and delivers the final value as a result.

 **Note:** All hidden layers usually use the same activation function. However, the output layer will typically use a different activation function from the hidden layers. The choice depends on the goal or type of prediction made by the model.

Feedforward vs. Backpropagation

When learning about neural networks, you will come across two essential terms describing the movement of information—feedforward and backpropagation.

Let's explore them.

 **Feedforward Propagation** - the flow of information occurs in the forward direction. The input is used to calculate some intermediate function in the hidden layer, which is then used to calculate an output.

In the feedforward propagation, the Activation Function is a mathematical “gate” in between the input feeding the current neuron and its output going to the next layer.

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



 **Backpropagation** - the weights of the network connections are repeatedly adjusted to minimize the difference between the actual output vector of the net and the desired output vector.

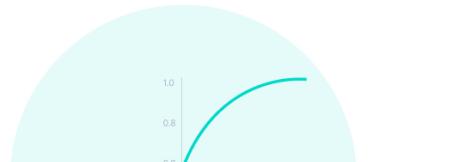
To put it simply—backpropagation aims to minimize the cost function by adjusting the network's weights and biases. The cost function gradients determine the level of adjustment with respect to parameters like activation function, weights, bias, etc.

Why do Neural Networks Need an Activation Function?

So we know what Activation Function is and what it does, but—

Why do Neural Networks need it?

Well, the purpose of an activation function is to add non-linearity to the neural network.



 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



Activation functions introduce an additional step at each layer during the forward propagation, but its computation is worth it. Here is why—

Let's suppose we have a neural network working *without* the activation functions.

In that case, every neuron will only be performing a linear transformation on the inputs using the weights and biases. It's because it doesn't matter how many hidden layers we attach in the neural network; all layers will behave in the same way because the composition of two linear functions is a linear function itself.

Although the neural network becomes simpler, learning any complex task is impossible, and our model would be just a linear regression model.

3 Types of Neural Networks Activation Functions

Now, as we've covered the essential concepts, let's go over the most popular neural networks activation functions.

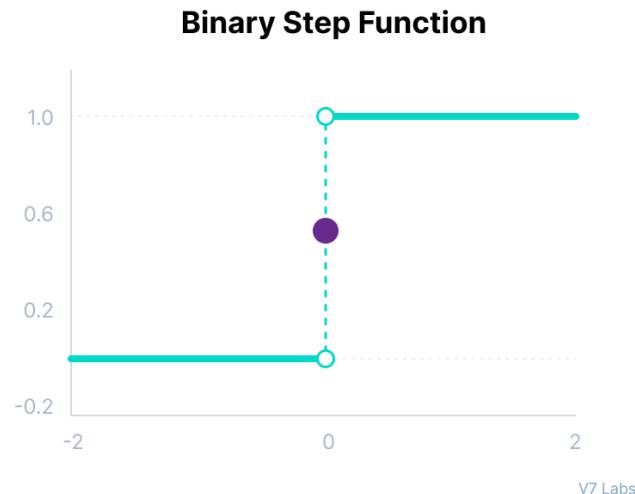
Binary Step Function

Binary step function depends on a threshold

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



else it is deactivated, meaning that its output is not passed on to the next hidden layer.



Binary Step Function

Mathematically it can be represented as:

Binary step

$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$$

Here are some of the limitations of binary step function:

- It cannot provide multi-value outputs—for example, it cannot be used for multi-class classification problems.
- The gradient of the step function is zero, which causes a hindrance in the backpropagation process.

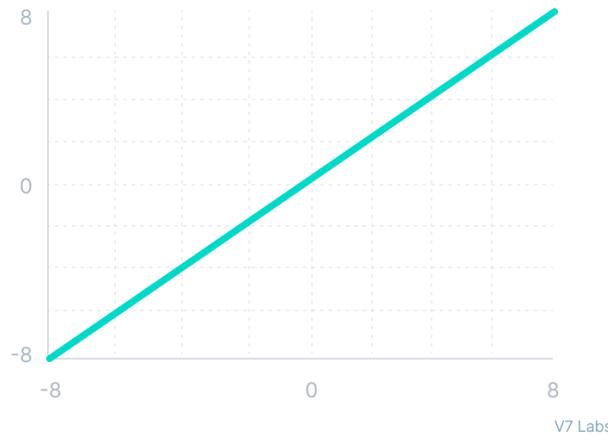
We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



(multiplied x1.0), is where the activation is proportional to the input.

The function doesn't do anything to the weighted sum of the input, it simply spits out the value it was given.

Linear Activation Function



Linear Activation Function

Mathematically it can be represented as:

Linear

$$f(x) = x$$

However, a linear activation function has two major problems :

- It's not possible to use backpropagation as the derivative of the function is a constant and has no relation to the input

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



first layer. So, essentially, a linear activation function turns the neural network into just one layer.

Non-Linear Activation Functions

The linear activation function shown above is simply a linear regression model.

Because of its limited power, this does not allow the model to create complex mappings between the network's inputs and outputs.

Non-linear activation functions solve the following limitations of linear activation functions:

- They allow backpropagation because now the derivative function would be related to the input, and it's possible to go back and understand which weights in the input neurons can provide a better prediction.
- They allow the stacking of multiple layers of neurons as the output would now be a non-linear combination of input passed through multiple layers. Any output can be represented as a functional computation in a neural network.

Now, let's have a look at ten different non-linear neural networks activation functions and their characteristics.

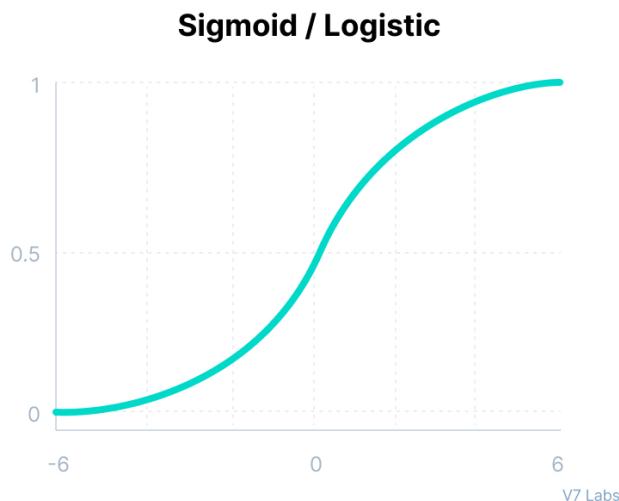
10 Non-Linear Neural Networks Activation

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



This function takes any real value as input and outputs values in the range of 0 to 1.

The larger the input (more positive), the closer the output value will be to 1.0, whereas the smaller the input (more negative), the closer the output will be to 0.0, as shown below.



Sigmoid/Logistic Activation Function

Mathematically it can be represented as:

Sigmoid / Logistic

$$f(x) = \frac{1}{1 + e^{-x}}$$

Here's why sigmoid/logistic activation function is one of the most widely used

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.

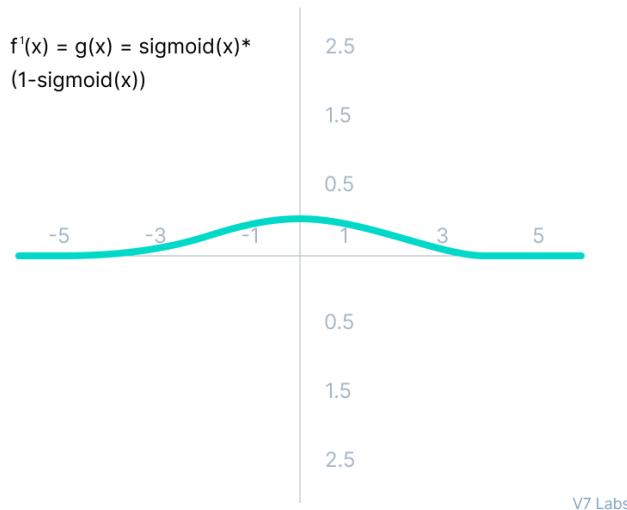


sigmoid is the right choice because of its range.

- The function is differentiable and provides a smooth gradient, i.e., preventing jumps in output values. This is represented by an S-shape of the sigmoid activation function.

The limitations of sigmoid function are discussed below:

- The derivative of the function is $f'(x) = \text{sigmoid}(x) * (1 - \text{sigmoid}(x))$.



The derivative of the Sigmoid Activation Function

As we can see from the above Figure, the gradient values are only significant for range -3 to 3, and the graph gets much flatter in other regions.

It implies that for values greater than 3 or less

than -3, the function will have very small

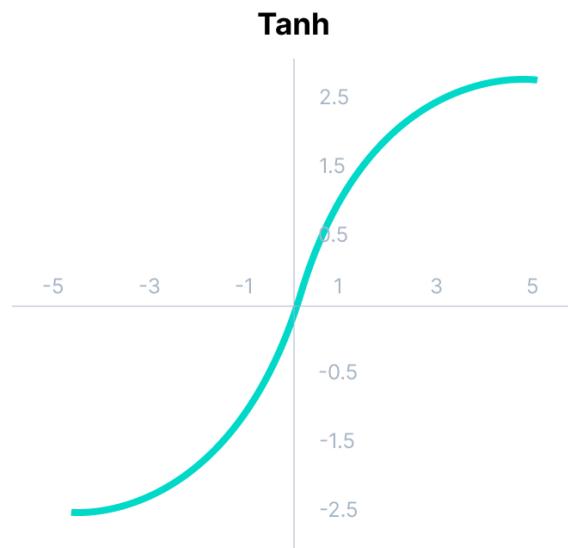
cookies We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



- The output of the logistic function is not symmetric around zero. So the output of all the neurons will be of the same sign. This makes the **training of the neural network** more difficult and unstable.

Tanh Function (Hyperbolic Tangent)

Tanh function is very similar to the sigmoid/logistic activation function, and even has the same S-shape with the difference in output range of -1 to 1. In Tanh, the larger the input (more positive), the closer the output value will be to 1.0, whereas the smaller the input (more negative), the closer the output will be to -1.0.



Tanh Function (Hyperbolic Tangent)

Mathematically it can be represented as:

$$\text{Tanh}$$

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.

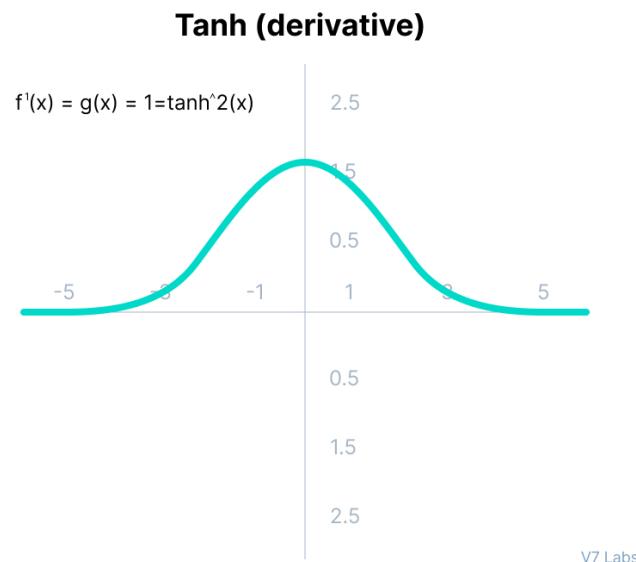


Advantages of using this activation function

are:

- The output of the tanh activation function is Zero centered; hence we can easily map the output values as strongly negative, neutral, or strongly positive.
- Usually used in hidden layers of a neural network as its values lie between -1 to; therefore, the mean for the hidden layer comes out to be 0 or very close to it. It helps in centering the data and makes learning for the next layer much easier.

Have a look at the gradient of the tanh activation function to understand its limitations.



Gradient of the Tanh Activation Function

As you can see—it also faces the problem of vanishing gradients similar to the sigmoid activation function. Plus the gradient of the

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



Note: Although both sigmoid and tanh face vanishing gradient issue, tanh is zero centered, and the gradients are not restricted to move in a certain direction. Therefore, in practice, tanh nonlinearity is always preferred to sigmoid nonlinearity.

ReLU Function

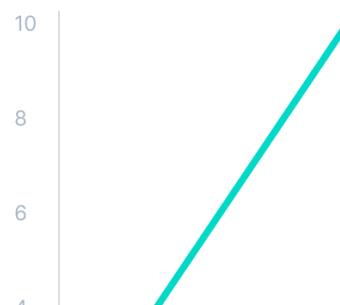
ReLU stands for Rectified Linear Unit.

Although it gives an impression of a linear function, ReLU has a derivative function and allows for backpropagation while simultaneously making it computationally efficient.

The main catch here is that the ReLU function does not activate all the neurons at the same time.

The neurons will only be deactivated if the output of the linear transformation is less than 0.

ReLU



We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



Mathematically it can be represented as:

ReLU

$$f(x) = \max(0, x)$$

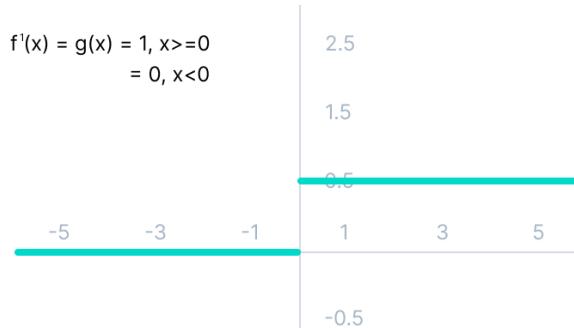
The advantages of using ReLU as an activation function are as follows:

- Since only a certain number of neurons are activated, the ReLU function is far more computationally efficient when compared to the sigmoid and tanh functions.
- ReLU accelerates the convergence of gradient descent towards the global minimum of the **loss function** due to its linear, non-saturating property.

The limitations faced by this function are:

- The Dying ReLU problem, which I explained below.

The Dying ReLU problem



We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



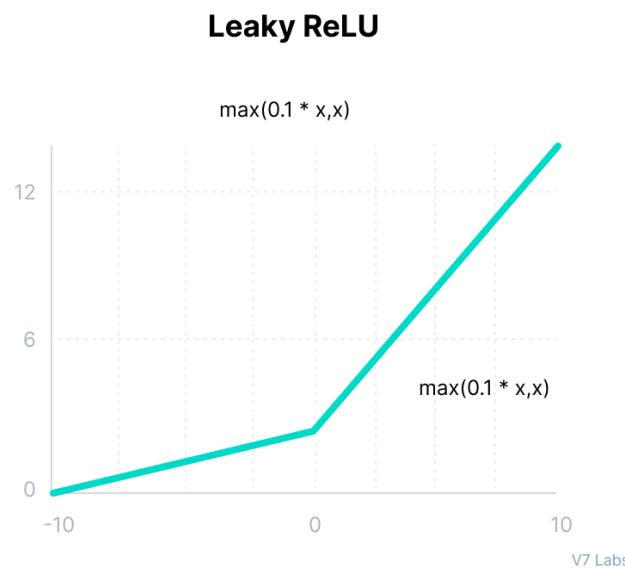
The negative side of the graph makes the gradient value zero. Due to this reason, during the backpropagation process, the weights and biases for some neurons are not updated. This can create dead neurons which never get activated.

- All the negative input values become zero immediately, which decreases the model's ability to fit or train from the data properly.

Note: For building the most reliable ML models, [split your data into train, validation, and test sets](#).

Leaky ReLU Function

Leaky ReLU is an improved version of ReLU function to solve the Dying ReLU problem as it has a small positive slope in the negative area.



We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



Leaky ReLU

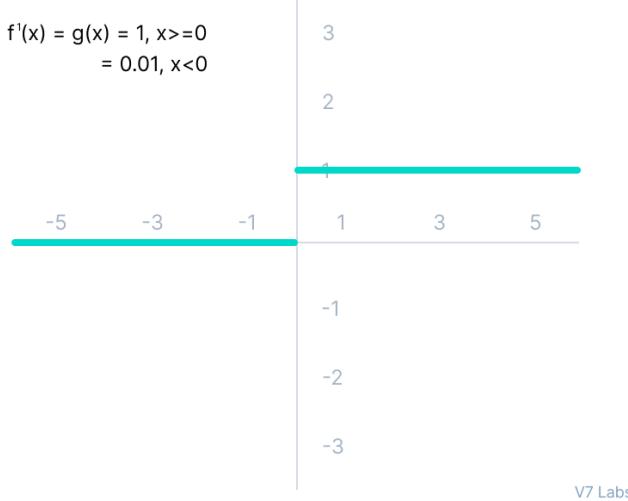
$$f(x) = \max(0.1x, x)$$

The advantages of Leaky ReLU are same as that of ReLU, in addition to the fact that it does enable backpropagation, even for negative input values.

By making this minor modification for negative input values, the gradient of the left side of the graph comes out to be a non-zero value. Therefore, we would no longer encounter dead neurons in that region.

Here is the derivative of the Leaky ReLU function.

Leaky ReLU (derivative)



The derivative of the Leaky ReLU function

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.

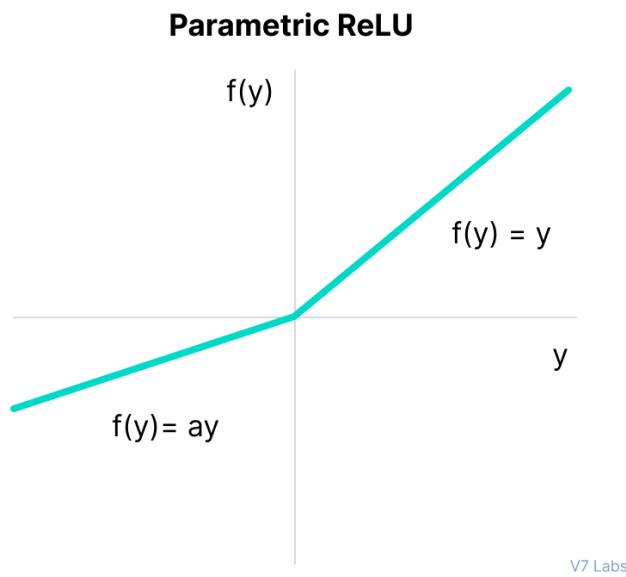


- The predictions may not be consistent for negative input values.
- The gradient for negative values is a small value that makes the learning of model parameters time-consuming.

Parametric ReLU Function

Parametric ReLU is another variant of ReLU that aims to solve the problem of gradient's becoming zero for the left half of the axis.

This function provides the slope of the negative part of the function as an argument
 a. By performing backpropagation, the most appropriate value of a is learnt.



Mathematically it can be represented as:

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



Where " a " is the slope parameter for negative values.

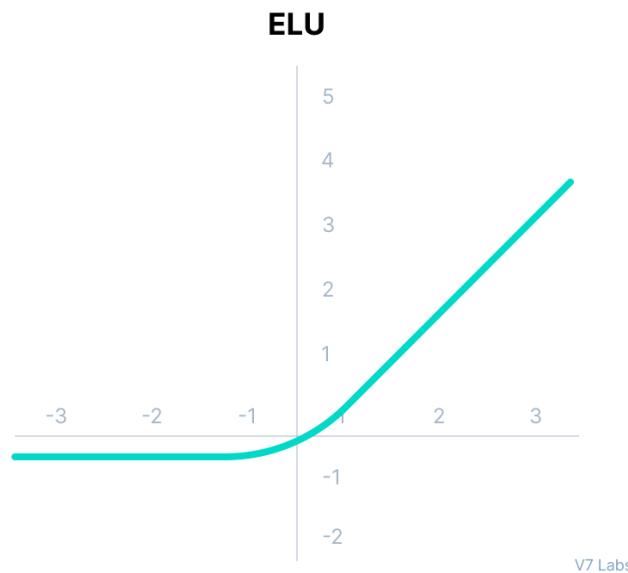
The parameterized ReLU function is used when the leaky ReLU function still fails at solving the problem of dead neurons, and the relevant information is not successfully passed to the next layer.

This function's limitation is that it may perform differently for different problems depending upon the value of slope parameter a .

Exponential Linear Units (ELUs) Function

Exponential Linear Unit, or ELU for short, is also a variant of ReLU that modifies the slope of the negative part of the function.

ELU uses a log curve to define the negative values unlike the leaky ReLU and Parametric ReLU functions with a straight line.



ELU Activation Function

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



ELU

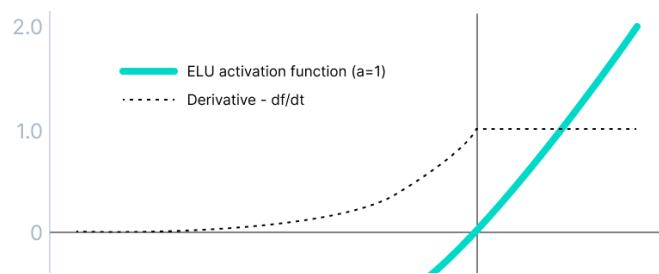
$$\begin{cases} x & \text{for } x \geq 0 \\ \alpha(e^x - 1) & \text{for } x < 0 \end{cases}$$

ELU is a strong alternative for f ReLU because of the following advantages:

- ELU becomes smooth slowly until its output equal to $-\alpha$ whereas RELU sharply smoothes.
- Avoids dead ReLU problem by introducing log curve for negative values of input. It helps the network nudge weights and biases in the right direction.

The limitations of the ELU function are as follow:

- It increases the computational time because of the exponential operation included
- No learning of the ' α ' value takes place
- Exploding gradient problem

ELU ($\alpha=1$) + Derivative

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



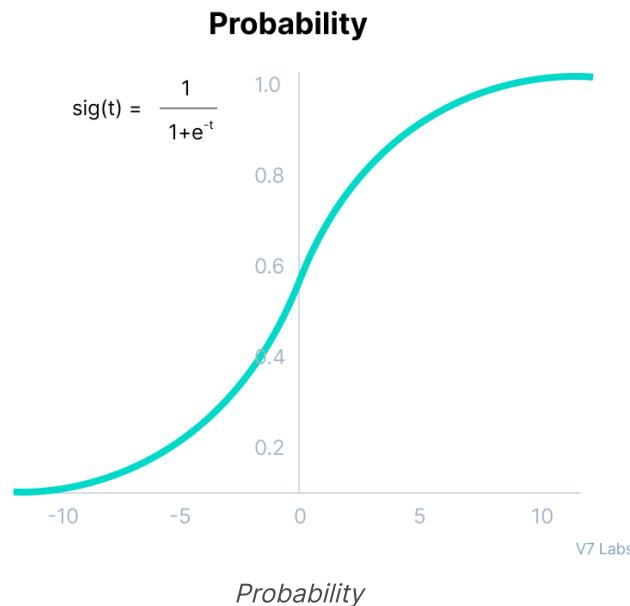
Mathematically it can be represented as:

ELU derivative

$$f'(x) = \begin{cases} 1 & \text{for } x \geq 0 \\ f(x) + \alpha & \text{for } x < 0 \end{cases}$$

Softmax Function

Before exploring the ins and outs of the Softmax activation function, we should focus on its building block—the sigmoid/logistic activation function that works on calculating probability values.



🍪 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



This function faces certain problems.

Let's suppose we have five output values of 0.8, 0.9, 0.7, 0.8, and 0.6, respectively. How can we move forward with it?

The answer is: We can't.

The above values don't make sense as the sum of all the classes/output probabilities should be equal to 1.

You see, the Softmax function is described as a combination of multiple sigmoids.

It calculates the relative probabilities. Similar to the sigmoid/logistic activation function, the SoftMax function returns the probability of each class.

It is most commonly used as an activation function for the last layer of the neural network in the case of multi-class classification.

Mathematically it can be represented as:

Softmax

$$\text{softmax}(z_i) = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

V7 Labs

Softmax Function

X

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.

Assume that you have three classes, meaning that there would be three neurons in the output layer. Now, suppose that your output from the neurons is [1.8, 0.9, 0.68].

Applying the softmax function over these values to give a probabilistic view will result in the following outcome: [0.58, 0.23, 0.19].

The function returns 1 for the largest probability index while it returns 0 for the other two array indexes. Here, giving full weight to index 0 and no weight to index 1 and index 2. So the output would be the class corresponding to the 1st neuron(index 0) out of three.

You can see now how softmax activation function make things easy for multi-class classification problems.

Swish

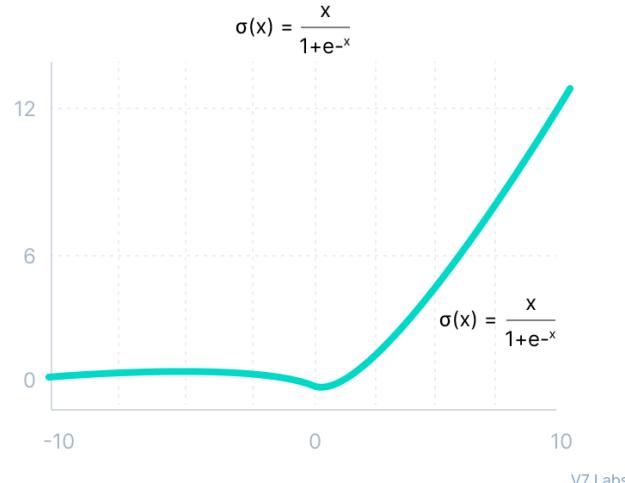
It is a self-gated activation function developed by researchers at Google.

Swish consistently matches or outperforms ReLU activation function on deep networks applied to various challenging domains such as [image classification](#), machine translation etc.

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



Swish



Swish Activation Function

This function is bounded below but unbounded above i.e. Y approaches to a constant value as X approaches negative infinity but Y approaches to infinity as X approaches infinity.

Mathematically it can be represented as:

Swish

$$f(x) = x * \text{sigmoid}(x)$$

Here are a few advantages of the Swish activation function over ReLU:

- Swish is a smooth function that means that it does not abruptly change

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



- Small negative values were zeroed out in ReLU activation function. However, those negative values may still be relevant for capturing patterns underlying the data. Large negative values are zeroed out for reasons of sparsity making it a win-win situation.
- The swish function being non-monotonous enhances the expression of input data and weight to be learnt.

Gaussian Error Linear Unit (GELU)

The Gaussian Error Linear Unit (GELU) activation function is compatible with BERT, ROBERTa, ALBERT, and other top NLP models. This activation function is motivated by combining properties from dropout, zoneout, and ReLUs.

ReLU and dropout together yield a neuron's output. ReLU does it deterministically by multiplying the input by zero or one (depending upon the input value being positive or negative) and dropout stochastically multiplying by zero.

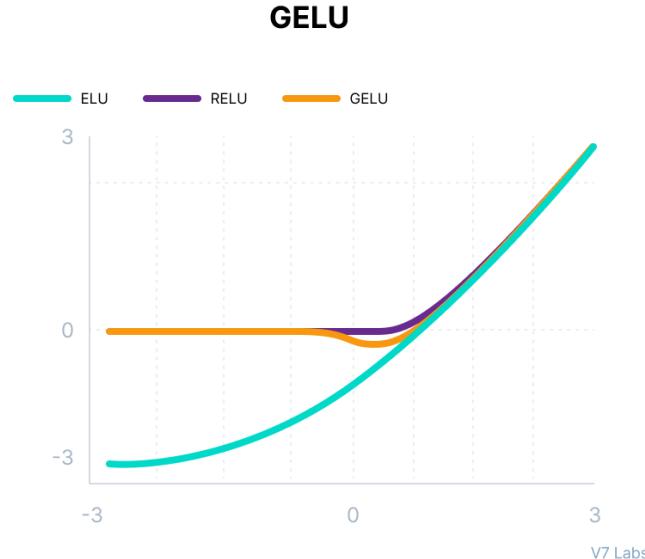
RNN regularizer called zoneout stochastically multiplies inputs by one.

We merge this functionality by multiplying the input by either zero or one which is stochastically determined and is dependent upon the input. We multiply the neuron input x by

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



especially with Batch Normalization.



Gaussian Error Linear Unit (GELU) Activation Function

Mathematically it can be represented as:

$$\begin{aligned} & \text{GELU} \\ & f(x) = xP(X \leq x) = x\Phi(x) \\ & = 0.5x \left(1 + \tanh \left[\sqrt{2/\pi} (x + 0.044715x^3) \right] \right) \end{aligned}$$

GELU nonlinearity is better than ReLU and ELU activations and finds performance improvements across all tasks in domains of **computer vision**, natural language processing, and speech recognition.

Scaled Exponential Linear Unit (SELU)

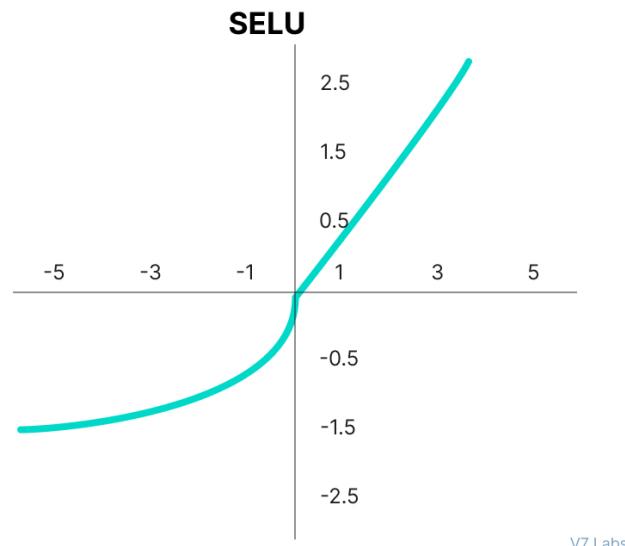
SELU was defined in self-normalizing networks and takes care of internal normalization which means each layer

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



SELU has both positive and negative values to shift the mean, which was impossible for ReLU activation function as it cannot output negative values.

Gradients can be used to adjust the variance. The activation function needs a region with a gradient larger than one to increase it.



SELU Activation Function

V7 Labs

Mathematically it can be represented as:

$$\text{SELU}$$

$$f(\alpha, x) = \lambda \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$

SELU has values of alpha α and lambda λ

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



network converges faster.

SELU is a relatively newer activation function and needs more papers on architectures such as CNNs and RNNs, where it is comparatively explored.

Why are deep neural networks hard to train?

There are two challenges you might encounter when training your deep neural networks.

Let's discuss them in more detail.

Vanishing Gradients

Like the sigmoid function, certain activation functions squish an ample input space into a small output space between 0 and 1.

Therefore, a large change in the input of the sigmoid function will cause a small change in the output. Hence, the derivative becomes small. For shallow networks with only a few layers that use these activations, this isn't a big problem.

However, when more layers are used, it can cause the gradient to be too small for training to work effectively.

Exploding Gradients

Exploding gradients are problems where significant error gradients accumulate and

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



The values of the weights can also become so large as to overflow and result in something called NaN values.

How to choose the right Activation Function?

You need to match your activation function for your output layer based on the type of prediction problem that you are solving—specifically, the type of predicted variable.

Here's what you should keep in mind.

As a rule of thumb, you can begin with using the ReLU activation function and then move over to other activation functions if ReLU doesn't provide optimum results.

And here are a few other guidelines to help you out.

1. ReLU activation function should only be used in the hidden layers.
2. Sigmoid/Logistic and Tanh functions should not be used in hidden layers as they make the model more susceptible to problems during training (due to vanishing gradients).
3. Swish function is used in neural networks having a depth greater than 40 layers.

Finally, a few rules for choosing the activation function for your output layer based on the type of prediction problem that you are

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



3. Multiclass Classification—Softmax

4. Multilabel Classification—Sigmoid

The activation function used in hidden layers is typically chosen based on the type of neural network architecture.

5. Convolutional Neural Network (CNN): ReLU activation function.

6. Recurrent Neural Network: Tanh and/or Sigmoid activation function.

And hey—use this cheatsheet to consolidate all the knowledge on the Neural Network Activation Functions that you've just acquired :)



We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



Neural Networks

Activation Functions in a Nutshell

Well done!

You've made it this far ;-) Now, let's have a quick recap of everything you've learnt in this tutorial:

- Activation Functions are used to introduce non-linearity in the network.
- A neural network will almost always have the same activation function in all hidden layers. This activation function should be differentiable so that the parameters of the network are learned in backpropagation.
- ReLU is the most commonly used activation function for hidden layers.
- While selecting an activation function, you must consider the problems it might face: vanishing and exploding gradients.
- Regarding the output layer, we must always consider the expected value range of the predictions. If it can be any numeric value (as in case of the regression problem) you can use the linear activation function or ReLU.
- Use Softmax or Sigmoid function for the

 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



40+ Data Science Interview Questions and Answers

[The Complete Guide to CVAT—Pros & Cons](#)

[5 Alternatives to Scale AI](#)

[YOLO: Real-Time Object Detection Explained](#)

[The Beginner's Guide to Contrastive Learning](#)

[9 Reinforcement Learning Real-Life Applications](#)

[Mean Average Precision \(mAP\) Explained: Everything You Need to Know](#)

[A Step-by-Step Guide to Text Annotation \[+Free OCR Tool\]](#)

[The Essential Guide to Data Augmentation in Deep Learning](#)

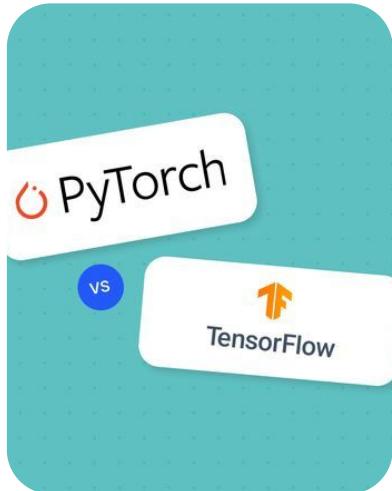


Pragati Baheti
Microsoft

Pragati is a software developer at Microsoft, and a deep learning enthusiast. She writes about the fundamental mathematics behind deep neural networks.

🍪 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.



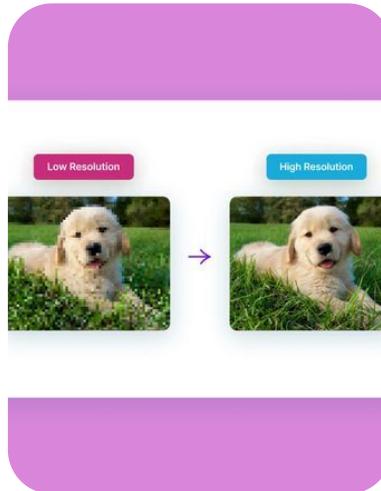


DEEP LEARNING

PyTorch vs TensorFlow: The Ultimate Decision Guide

Konstantinos Poulinakis

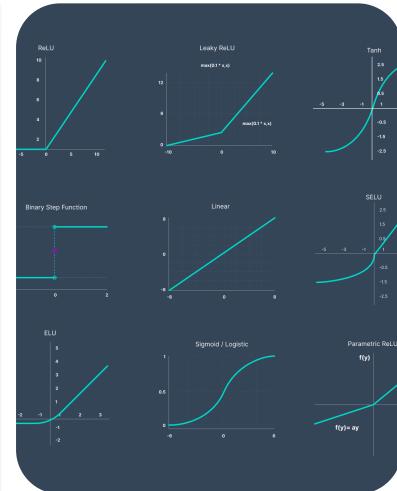
15 min read



DEEP LEARNING

Deep Learning for Image Super-Resolution [incl. Architectures]

Rohit Kundu 13 min read



DEEP LEARNING

Activation Functions in Neural Networks [12 Types & Use Cases]

Pragati Baheti 14 min read

Gain control of your training data

15,000+ ML engineers can't be wrong

Your email

Request a demo

We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.





COMPANY	PLATFORM	RESOURCES	INDUSTRIES	COMPARE
About	Image Annotation	Blog	Agriculture	V7 vs Scale AI
Pricing	Auto Annotation	Product Updates	Automotive	V7 vs Superannotate
Contact Us	Video Annotation	Playbooks	Construction	V7 vs Labelbox
Jobs	Dataset Management	Webinars	Energy	V7 vs Roboflow
News	Document Processing	Documentation	Food & Beverage	V7 vs Dataloop
Partners	Model Management	Academy	Healthcare	V7 vs Supervisely
Data Security	Workflows	Open Datasets	Insurance & Finance	V7 vs Encord
	Labeling Services	Community	Life Sciences & Biotech	V7 vs CVAT
		ML Glossary	Logistics	
		Ethics & CoC	Manufacturing	
			Retail	
			Software & Internet	
			Sports	

Subscribe to our
monthly newsletter

Enter yo! →

News, feature
releases, and blog
articles on AI

©V7Labs · Terms & Privacy



🍪 We use cookies on our website to keep our website safe, improve your experience, and for marketing. We won't turn them on until you accept or you can adjust your preferences.

